

Neural Networks I

Neural networks are universal function approximators that can approximate any continuous function provided the network has enough neurons (nodes).

- True
- False

Which term most precisely describes the process that process allows you to calculate the output of a neural network assuming you know the model weights?

- Backpropagation
- Bias nodes
- Forward propagation
- Gradient descent
- Hidden units
- Neurons
- Output nodes

Which term most precisely describes the process that allows you to calculate the gradient of a neural network model's cost function with respect to the weights?

- Backpropagation
- Bias nodes
- Forward propagation
- Gradient descent
- Hidden units
- Neurons
- Output nodes

Which term most precisely describes the optimization method used to iteratively converge upon the model weights that best minimize the cost function of the neural network?

- Backpropagation
- Bias nodes
- Forward propagation
- Gradient descent
- Hidden units
- Neurons
- Output nodes

What's the hype around neural networks?


Character/handwriting recognition


Language translation


Medical diagnosis



Automated financial trading systems

And some other interesting computer vision applications...

Search Aerial Imagery (NAIP) 



Top 1000 Closest Matches:  [Clear Results](#)



Geovisual Search


<https://search.descarteslabs.com/>

Image-to-image translation

TOOL

- background
- wall
- door
- window**
- window sill
- window head
- shutter
- balcony
- trim
- cornice
- column
- entrance

INPUT



pix2pix
process

OUTPUT




Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

Image-to-image translation


Labels to Street Scene



Labels to Facade



BW to Color



Aerial to Map



Day to Night



Edges to Photo




Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

Image-to-image translation



Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." arXiv preprint (2017).

Image Style Transfer



Dumoulin, Vincent, Jonathon Shlens, and Manjunath Kudlur. "A learned representation for artistic style." CoRR, abs/1610.07629 2.4 (2016): 5.

What makes neural networks special?

Neural network learning is representation learning

Previous ML algorithms we discussed required us to manually determine feature transformations

Neural networks **learn** feature transformations

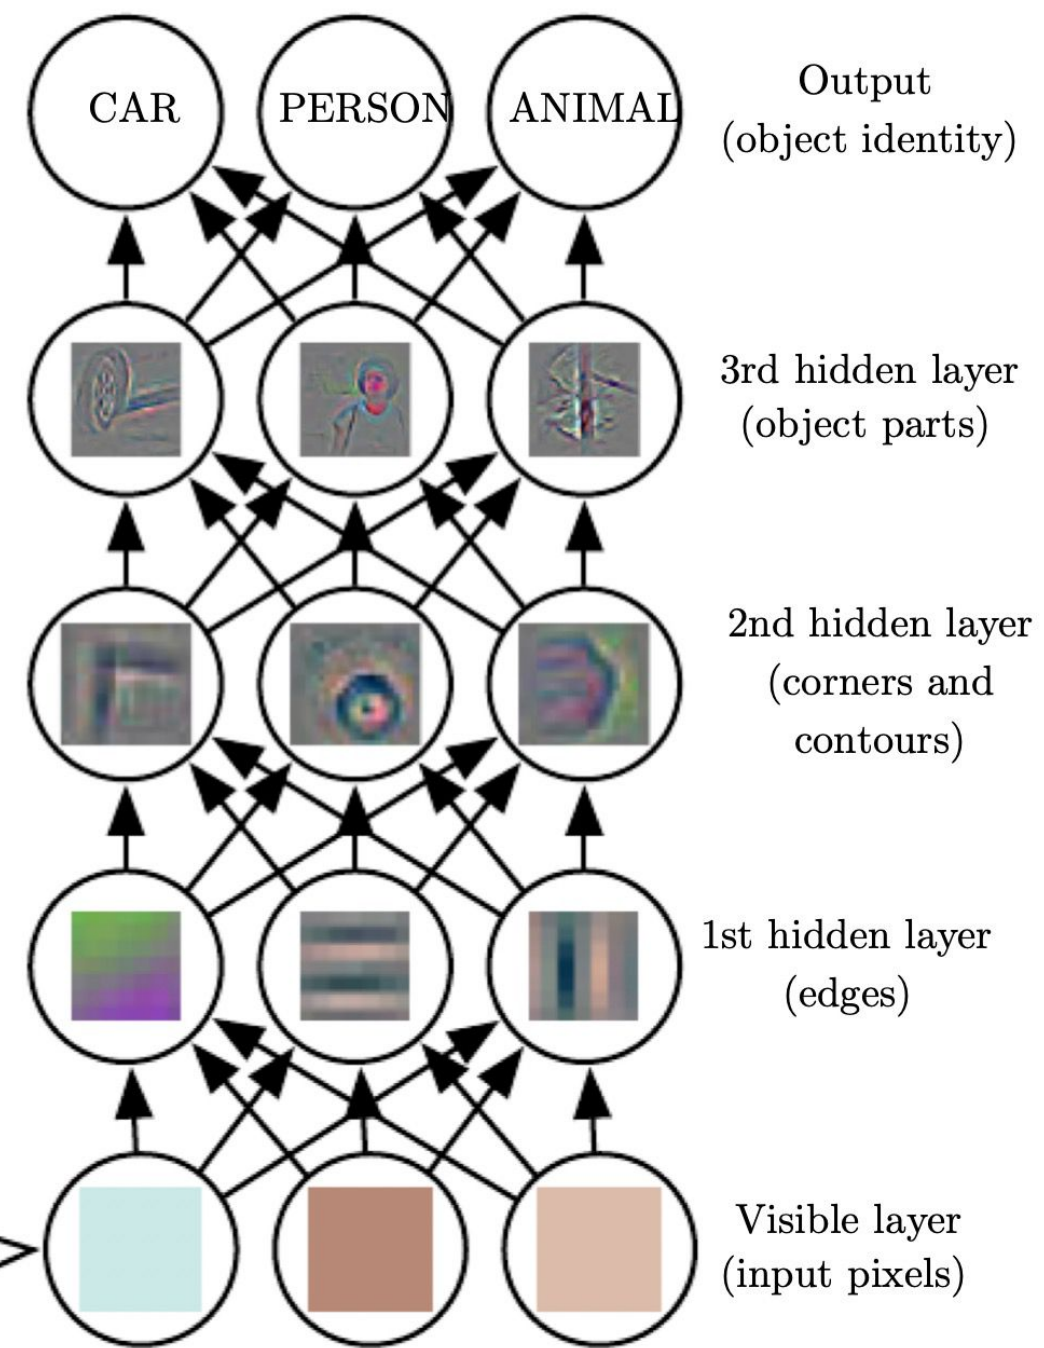



Image from Goodfellow, I., Bengio, Y., Courville, A. and Bengio, Y., 2016.
Deep learning (Vol. 1, No. 2). Cambridge: MIT press.

What is a neural network and **how does it work?**

How do we **optimize model weights?**
(i.e. how do we fit our model to data)


What are the challenges of using neural networks?

Recall our goal in supervised learning



Perceptron


$$\hat{f}(x) = \text{sign} \left(\sum_{i=0}^p w_i x_i \right)$$




Source: Abu-Mostafa, Learning from Data, Caltech

Multilayer Perceptron

What if we stuck multiple perceptrons together?



Perceptron #1




$$\hat{f}_1(\mathbf{x}) = \text{sign}(\mathbf{w}_1^T \mathbf{x})$$

The sharp
boundary is due to
our sign function



Perceptron #2




$$\hat{f}_2(\mathbf{x}) = \text{sign}(\mathbf{w}_2^T \mathbf{x})$$

Source: Abu-Mostafa, Learning from Data, Caltech

Multilayer perceptron:


$$\hat{f}(\mathbf{x}) = \begin{cases} +1 & \hat{f}_1(\mathbf{x}) \neq \hat{f}_2(\mathbf{x}) \\ -1 & \hat{f}_1(\mathbf{x}) = \hat{f}_2(\mathbf{x}) \end{cases}$$

Perceptron #1




$$\hat{f}_1(\mathbf{x}) = \text{sign}(\mathbf{w}_1^T \mathbf{x})$$

Perceptron #2




$$\hat{f}_2(\mathbf{x}) = \text{sign}(\mathbf{w}_2^T \mathbf{x})$$




Source: Abu-Mostafa, Learning from Data, Caltech


Multilayer Perceptron



Target



8 perceptrons



16 perceptrons

The more nodes/neurons, the more flexible is the model


Source: Abu-Mostafa, Learning from Data, Caltech

Universal function approximation

“A **feedforward network** with a single layer is sufficient to represent **any function**, but the layer may be infeasibly large and may fail to learn and generalize correctly.”

Ian Goodfellow, Deep Learning
Creator of generative adversarial networks

Input nodes / neurons



Simply passes the input value to the next layer

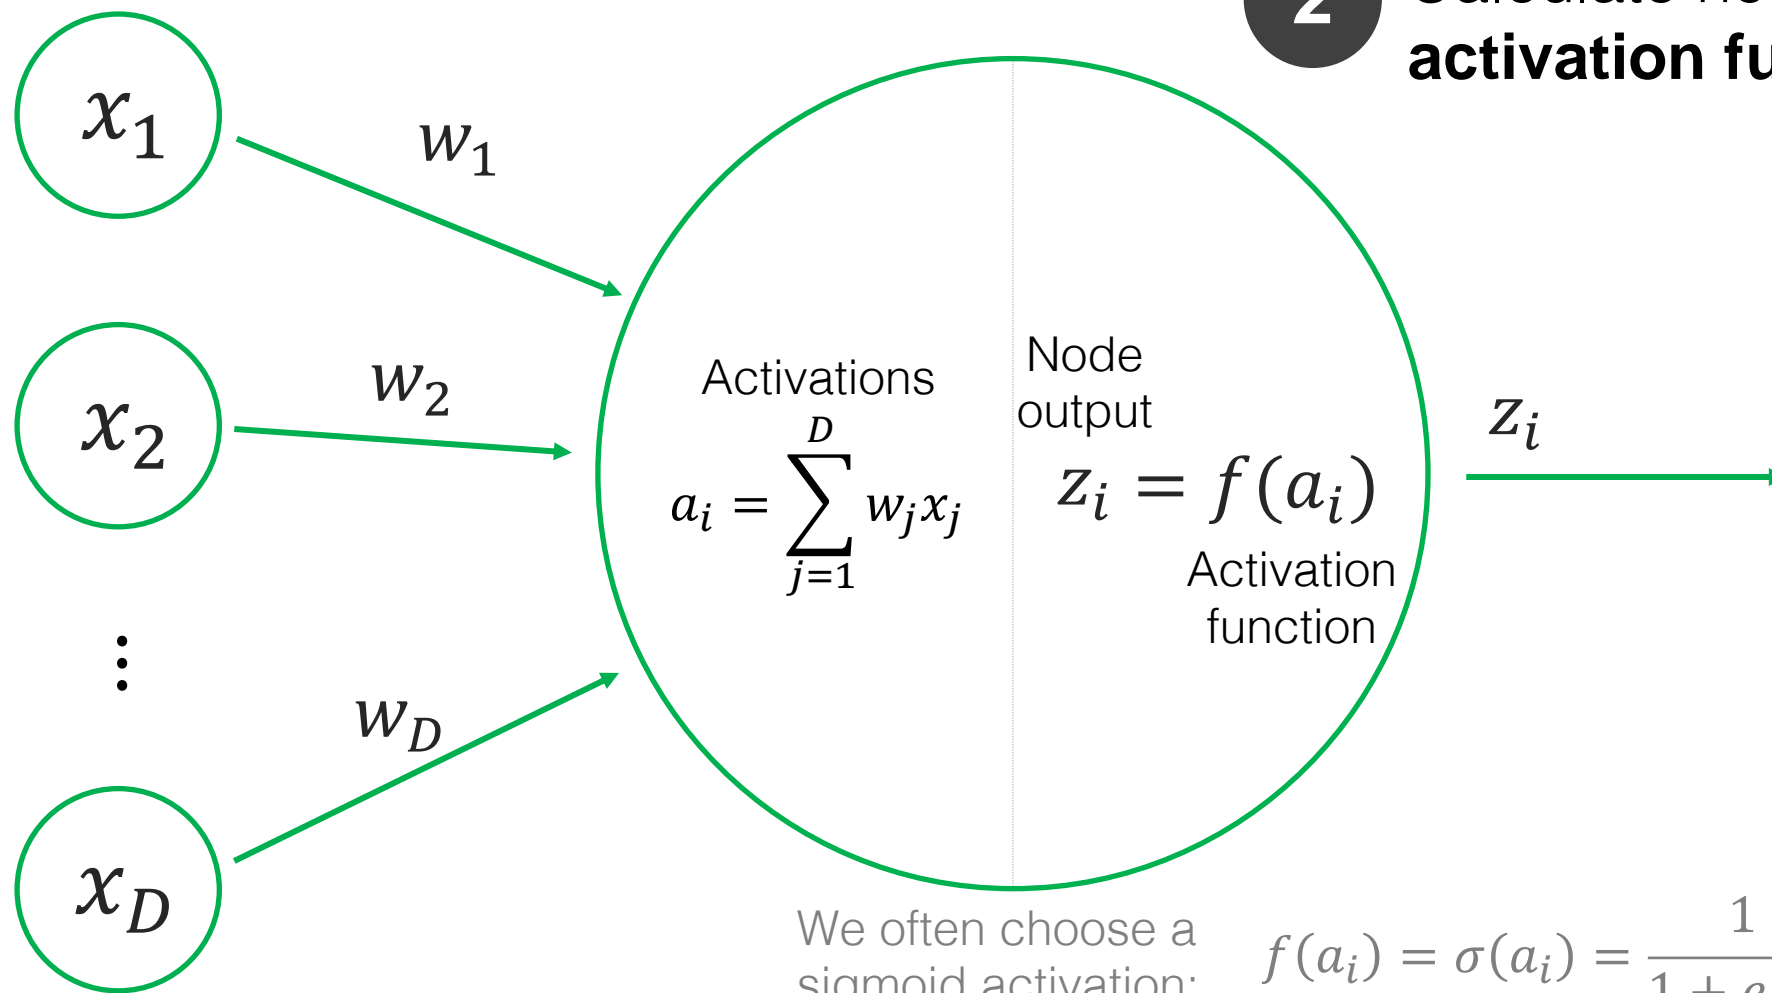
Hidden & output nodes

1


Calculate the **activations**: linear combinations of weights and the last layer's output

2


Calculate node output: apply the **activation function** to the activations



Represented as:



Simple Neural Network



Notational shorthand:
(a more precise
alternative notation)

$$w_{ij} = w_{ij}^{(1)}$$

$$z_i = z_i^{(1)}$$

$$w_{ij} = w_{ij}^{(2)}$$

$$y = z_1^{(2)}$$

$w_{ij}^{(k)}$

- Layer k
- From node j
(in the last layer)
- to node i
(in the next layer)

Forward Propagation


Calculating the output from input

$$a_1 = (0.9)(1.0) + (0.3)(0.5) = 1.05 \quad \text{Hidden layer calculations}$$

$$a_2 = (0.2)(1.0) + (0.8)(0.5) = 0.6$$

$$z_1 = \sigma(a_1) = \sigma(1.05) = 0.741$$

$$z_2 = \sigma(a_2) = \sigma(0.6) = 0.646$$



Output layer calculations

$$a_y = (0.2)(0.741) + (0.5)(0.646) = 0.471$$

$$y = a_y = 0.471 \quad \text{Regression}$$


Alternatively...

$$y = \sigma(a_y) = \sigma(0.471) = 0.616 \quad \text{Classification}$$

$$\sigma(a_i) = \frac{1}{1 + e^{-a_i}}$$

Forward Propagation

Calculating the output from input



Hidden layer matrix calculations

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \end{bmatrix} \quad \mathbf{z} = \begin{bmatrix} z_1 \\ z_2 \end{bmatrix}$$

$$\mathbf{W} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \begin{array}{l} \longrightarrow \text{The weights INTO node } z_1 \\ \longrightarrow \text{The weights INTO node } z_2 \end{array}$$


$$\mathbf{a} = \mathbf{W}\mathbf{x} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$= \begin{bmatrix} w_{11}x_1 + w_{12}x_2 \\ w_{21}x_1 + w_{22}x_2 \end{bmatrix}$$

$$\mathbf{z} = \sigma(\mathbf{a}) = \begin{bmatrix} \sigma(w_{11}x_1 + w_{12}x_2) \\ \sigma(w_{21}x_1 + w_{22}x_2) \end{bmatrix}$$

Forward Propagation

Example neural network with $L = 3$ layers and the i th layer has N_i nodes



Simple steps for forward propagation:

$$\text{For } i = 1 \text{ to } L - 1:$$
$$\mathbf{z}^{(i)} = \sigma(\mathbf{W}^{(i)} \mathbf{z}^{(i-1)})$$

Where:

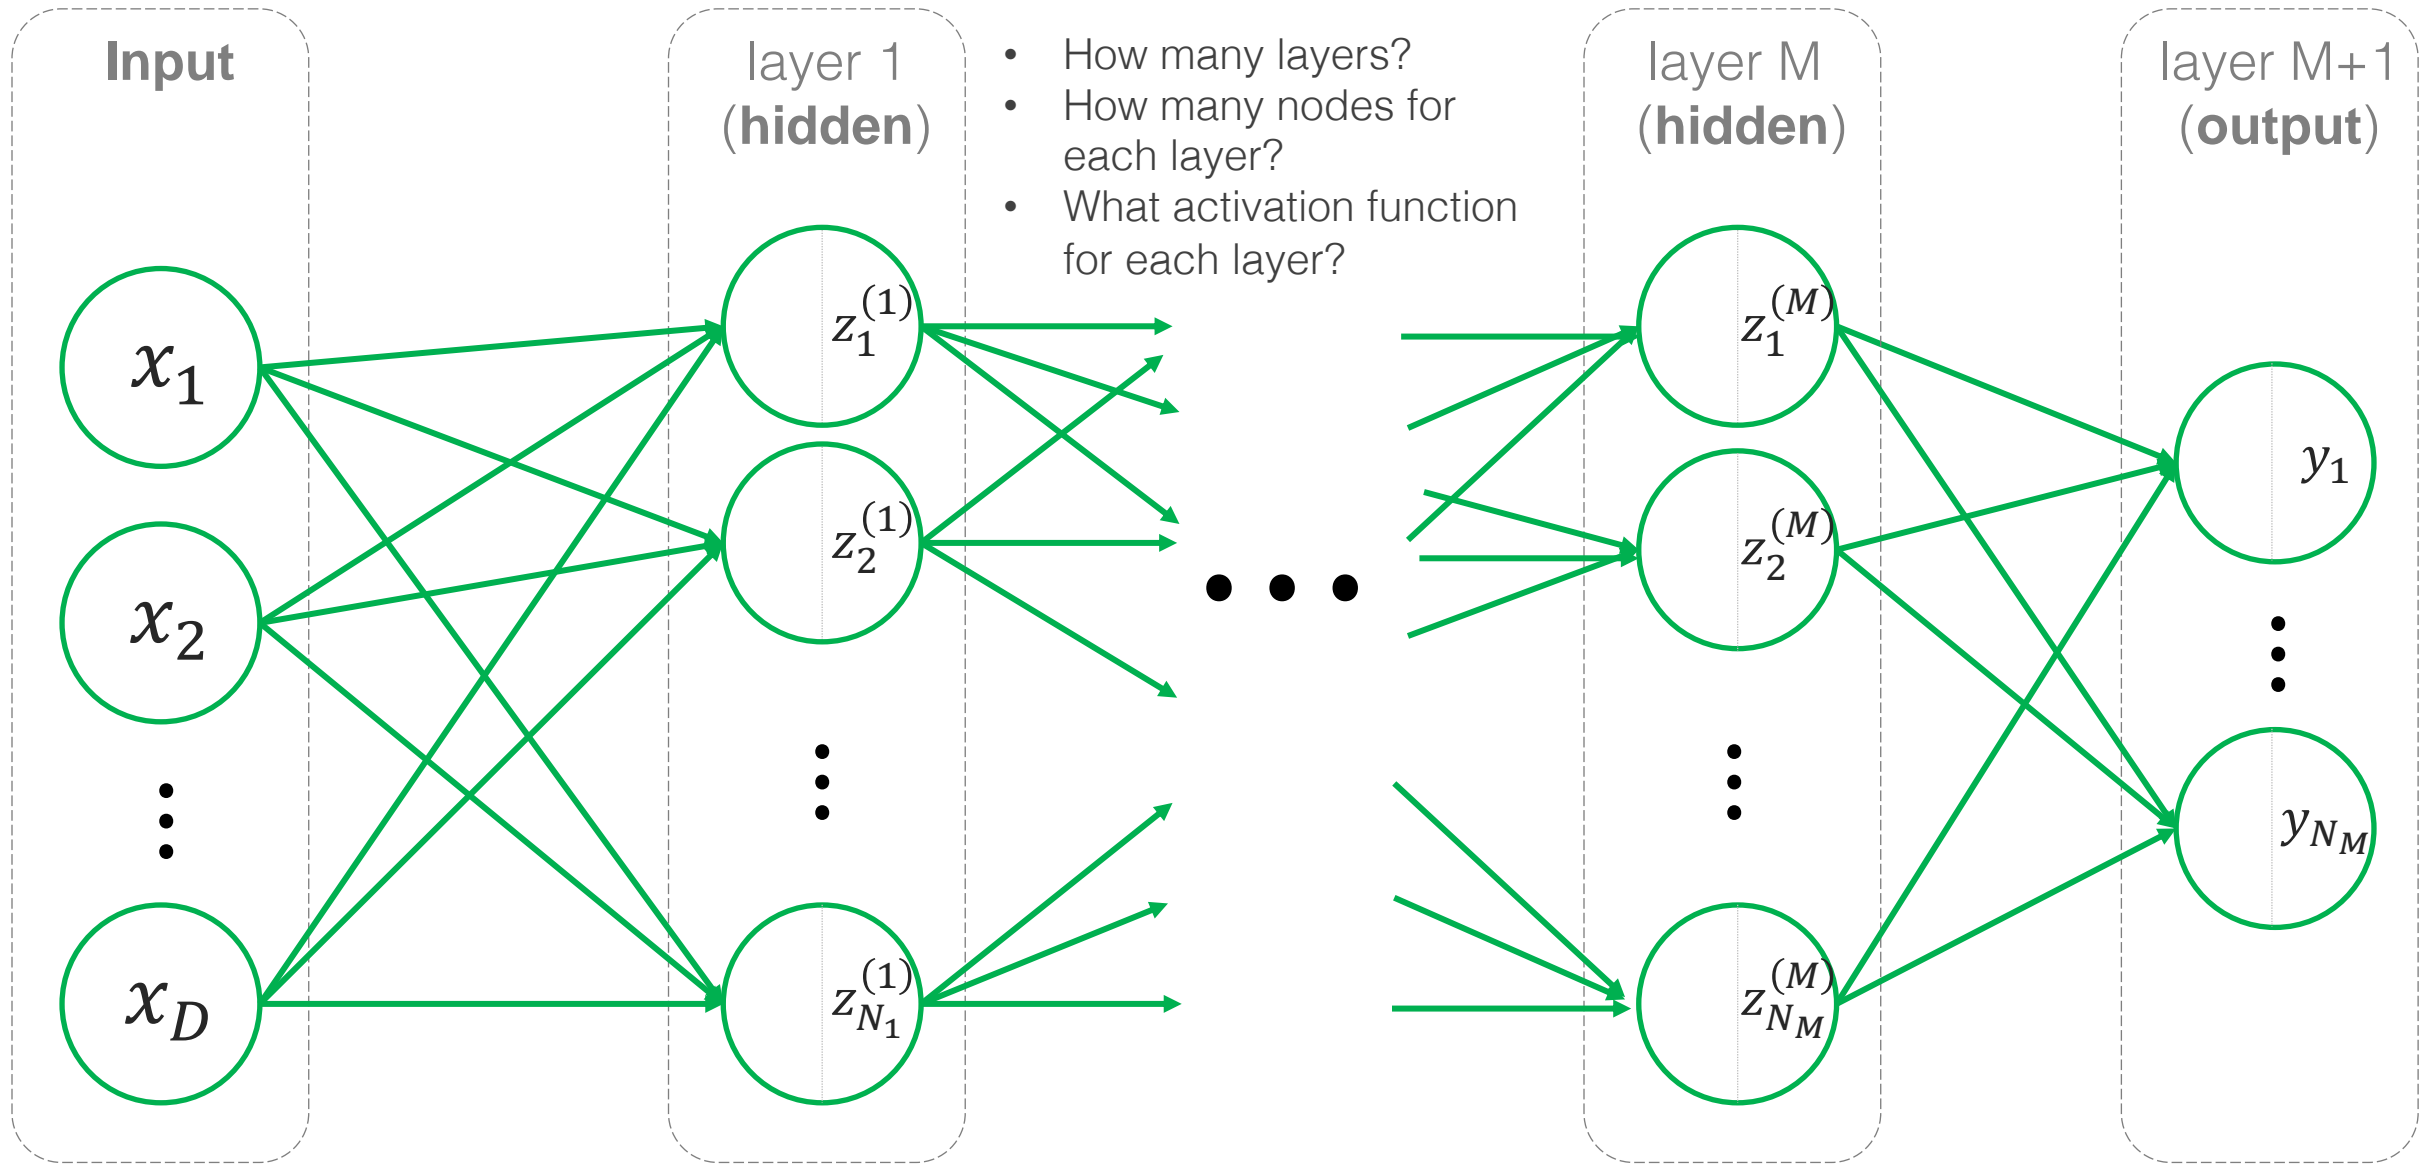
$$\mathbf{z}^{(0)} = \mathbf{x}$$

$$\hat{\mathbf{y}} = \mathbf{z}^{(L)}$$

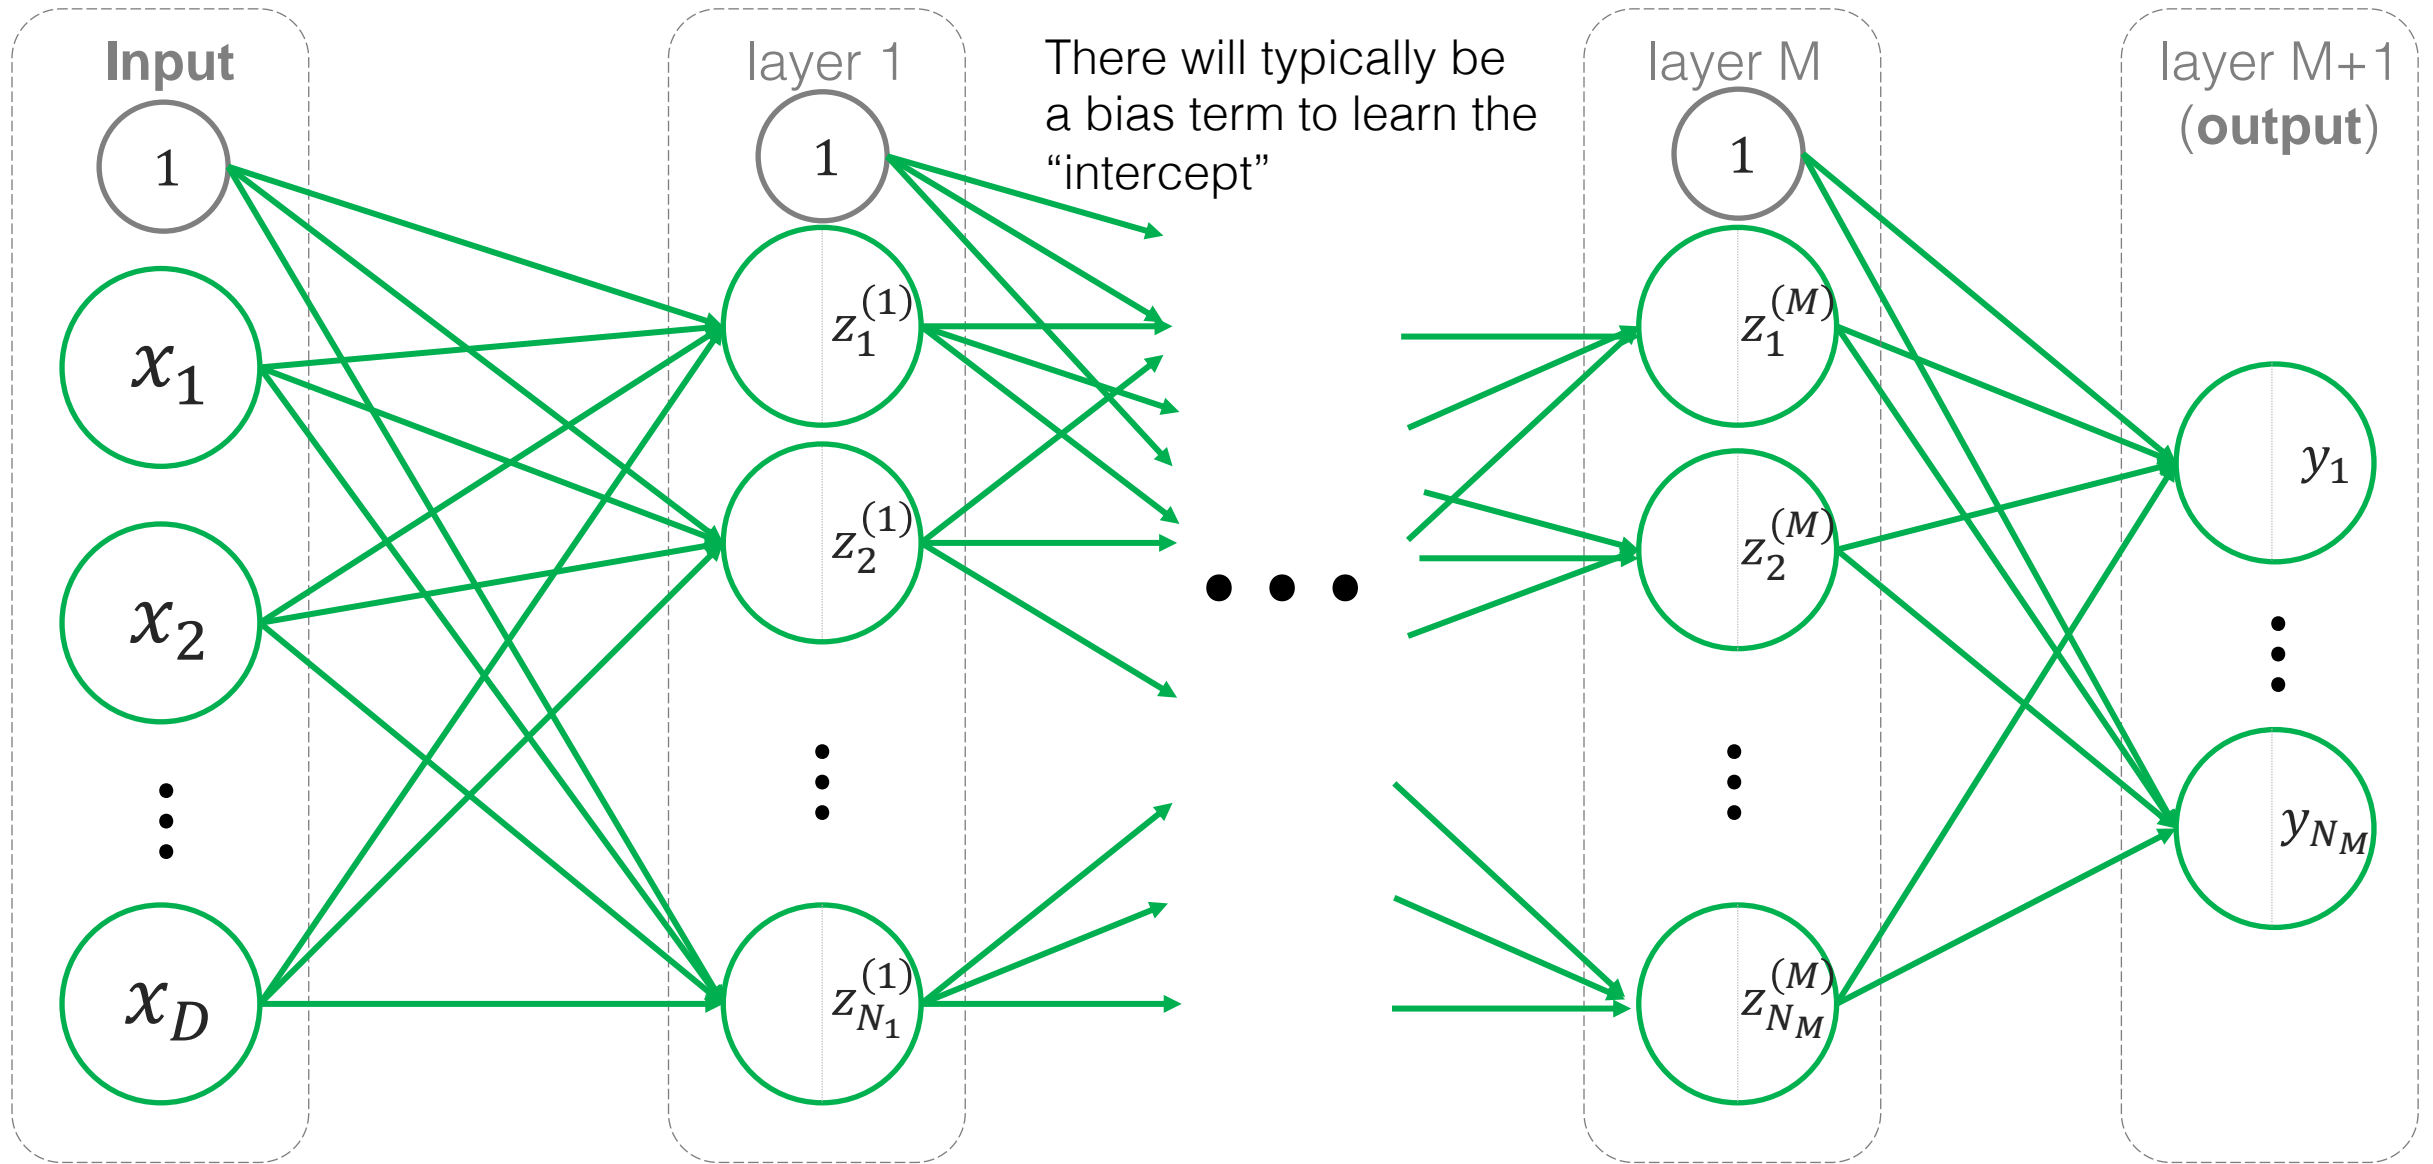
Prediction error is measured:

$$E_n = \frac{1}{2} (\hat{y}_n - y_n)^2$$

Neural networks can be customized



Neural networks can be customized














Input

Hidden 1


Hidden 2

Output



Multilayer neural nets for **regression** can build up from basic building blocks to more complex structures

From binary to multiclass classification



For **binary classification** with a sigmoid activation function, the output is between zero and one, so threshold this value to assign the class

Softmax

Generalization of the logistic function to multiple dimensions

Output activations

$$\begin{bmatrix} 5.1 \\ 4.2 \\ -3.1 \\ 0.7 \end{bmatrix}$$

Activation Function


$$\text{Softmax} \quad \frac{e^{a_i}}{\sum_{n=1}^M e^{a_n}}$$

Output

$$\begin{bmatrix} 0.7046 \\ 0.2865 \\ 0.0002 \\ 0.0087 \end{bmatrix}$$

Always sums to 1
(normalizes to be a
probability distribution)

From binary to multiclass classification



For **multiclass problems**, we can have multiple outputs and use a softmax function:

$$y_i = g(a_i) = \frac{e^{a_i}}{\sum_{n=1}^M e^{a_n}}$$

Choose the largest y value as the predicted class

As with many aspects of neural networks this is one of a number of approaches

Next time...

What is a neural network and **how does it work?**

How do we **optimize model weights?**

(i.e. how do we fit our model to data)

What are the challenges of using neural networks?