

# Reinforcement Learning III

# Final Report

- Header
- Abstract
- Introduction
- Background
- Data
- Methods
- Results
- Conclusions
- Roles
- Timeline of Activity
- References
- 2,500 words max (excluding references, roles, timeline of activity)
- Complete grading rubric will be posted this week

# Other Final Project Deliverables

- Video
  - Content is key – make sure you tell a clear story
  - Is it engaging and make use of the visual medium? Great opportunity to use images
  - Up to 5 min (NO LONGER)
- Github Repo
  - Contains a descriptive README.md file that explains what the repo is for, and how to use the code to reproduce your work (including how to set it up to run)
  - Is well commented throughout all files
  - Lists all dependencies in a requirements.txt file
  - Informs the user how to get the data and includes all preprocessing code
  - It actually runs and does what it says
- Peer Evaluation (criteria on the website)

## Do...

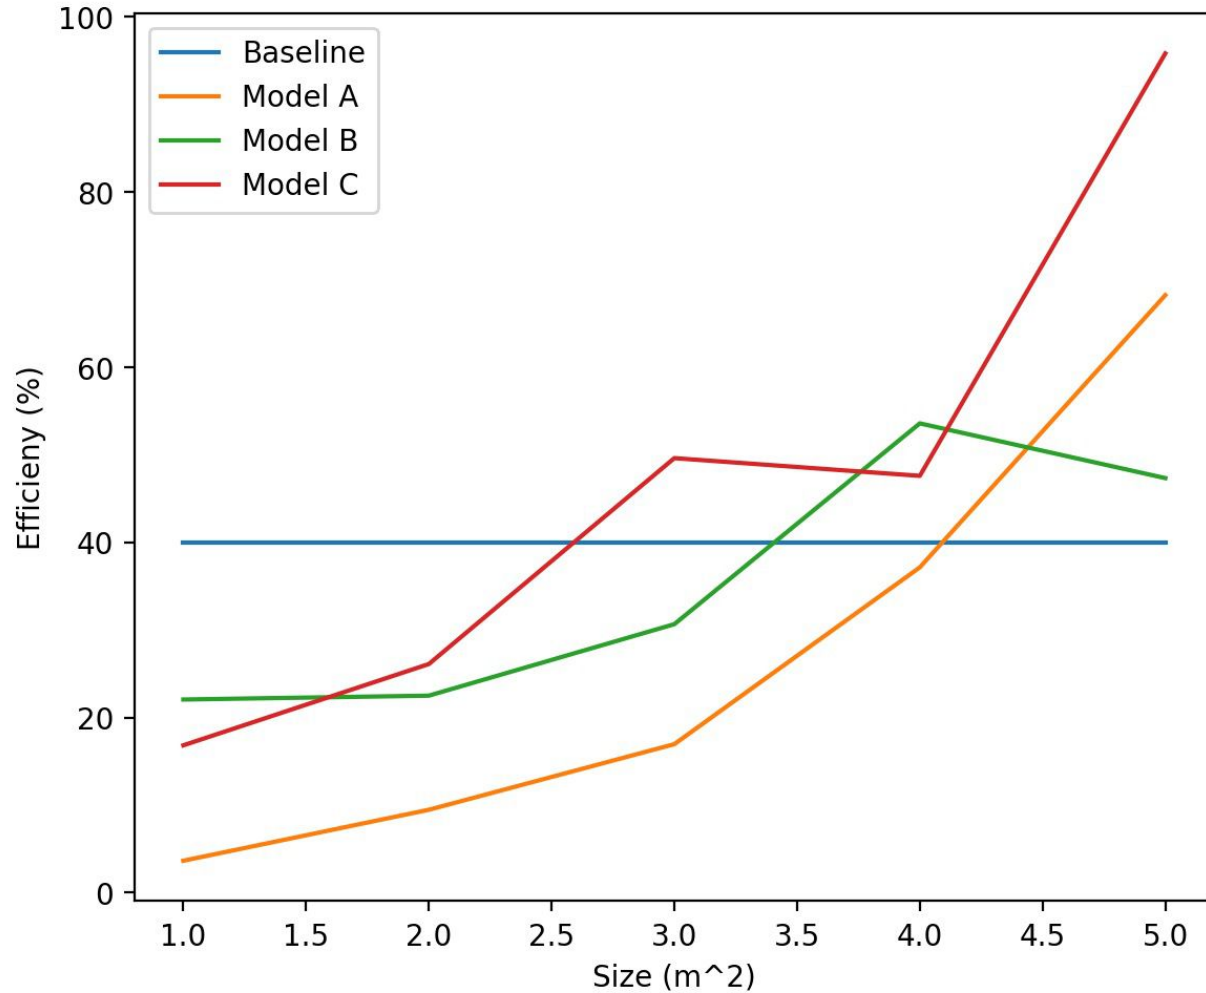
1. If you can show it as a figure, **do it**
2. Tell a story, don't write a logbook
3. Avoid unnecessary equations
4. Make your results/findings/conclusions clear and relate them back to your motivating problem/question. Avoid overly vague conclusions.
5. Make sure every figure has a contribution
6. Use flowcharts / diagrams to explain processes
7. Write professionally: avoid colloquialisms, check for spelling, grammar, punctuation, etc. Also check figure caption text
8. Reference your figures in the paper (e.g. Figure 1....). If a figure is not referenced, it doesn't have a point
9. Be precise in your language. For example: what do you mean when you say "better" – better how and in comparison to what?

## Don't...


1. Do **not** use screenshots, code snippets, or variable names from code in the report or video. Use figures or tables for code output
2. Do not use random pictures of ML model architectures unless it makes a point in your report

# Make your content EASY to understand


Satisfactory



Exceptional



# Agent-environment Interaction



**Agent** at each step  $t$ ...

Encounters state,  $s_t$   
Executes action  $a_t$   
Receives scalar reward,  $r_{t+1}$

**Environment** at each step  $t$ ...

Receives action  $a_t$   
Transitions to state,  $s_{t+1}$   
Emits scalar reward,  $r_{t+1}$


**Actions:** choices made by the agent

**States:** basis on which choices are made

**Rewards:** define the agent's goals

David Silver, 2015

# Reinforcement Learning Components



**Policy** (agent behavior),  $\pi(s)$

**Reward function** (the goal),  $r_t$

**Value functions** (expected returns),  
 $v(s)$  State value  
 $q(s, a)$  Action value

# Markov Models

States are  
**Fully Observable**

States are  
**Partially Observable**

**Autonomous**  
(no actions;  
make predictions)

Markov Chain,  
Markov Reward Process

Hidden Markov Model  
(HMM)

**Controlled**  
(can take actions)

Markov Decision  
Process (MDP)

Partially Observable  
Markov Decision  
Process (POMDP)

## Applications

HMMs: time series ML, e.g. speech + handwriting recognition, bioinformatics

MDPs: used extensively for reinforcement learning



# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)

**MDPs form the basis for most reinforcement learning environments**

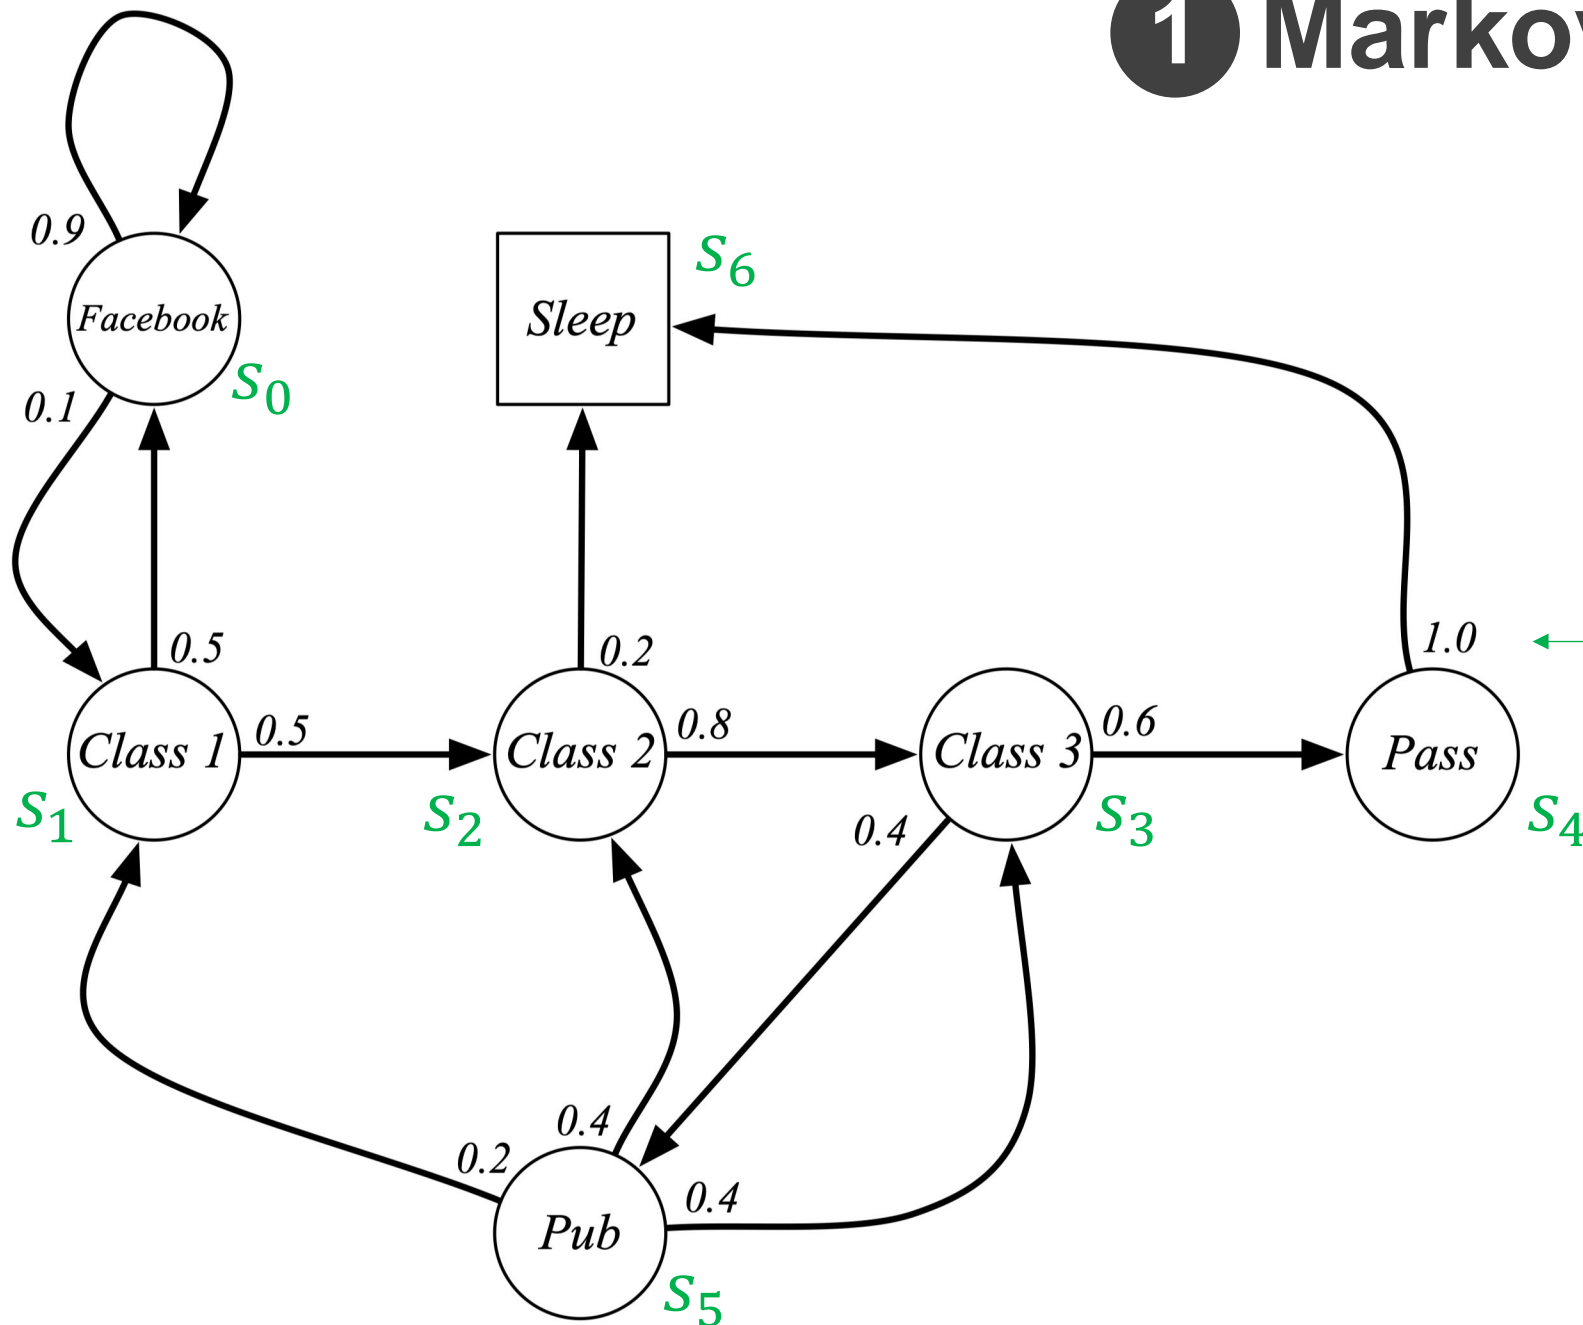
Adapted from David Silver, 2015

# 1 Markov Chain Example

Components:

State space  $S$ ,

Transition probabilities  $P$



$$P_{46} = P_{ss'}$$

Sample Episodes:

C1,C2,Sleep

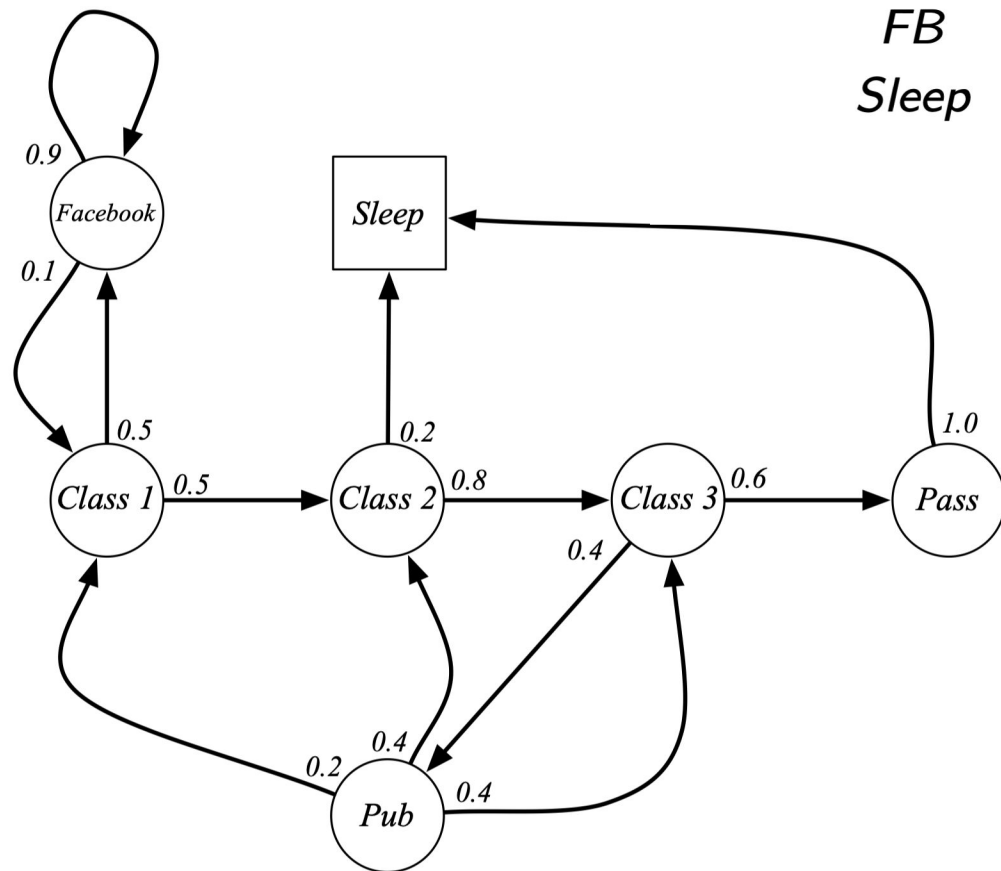
C1,FB,FB,FB,C1,C2,C3,Pass,Sleep

Example from David Silver, UCL, 2015

# Markov Chain

$$\mathcal{P} = \begin{matrix} & \begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix} \end{matrix}$$

	C1	C2	C3	Pass	Pub	FB	Sleep
C1		0.5				0.5	
C2			0.8				0.2
C3				0.6	0.4		
Pass	0.2	0.4	0.4				1.0
Pub	0.1					0.9	
FB							
Sleep							1



State transition probability matrix,  $P_{ss'}$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process


### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



Recall that returns, let's call  $G_t$ , are the total discounted rewards from time  $t$ :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process


### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process


### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0.9$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015

# “Backup” property of state value functions

$$v(s_t) = E[G_t | S = s_t] \quad \text{where } G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$$

$$= E[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots | S = s_t]$$

$$= E[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} \dots) | S = s_t]$$

$$= E[R_{t+1} + \gamma G_{t+1} | S = s_t]$$


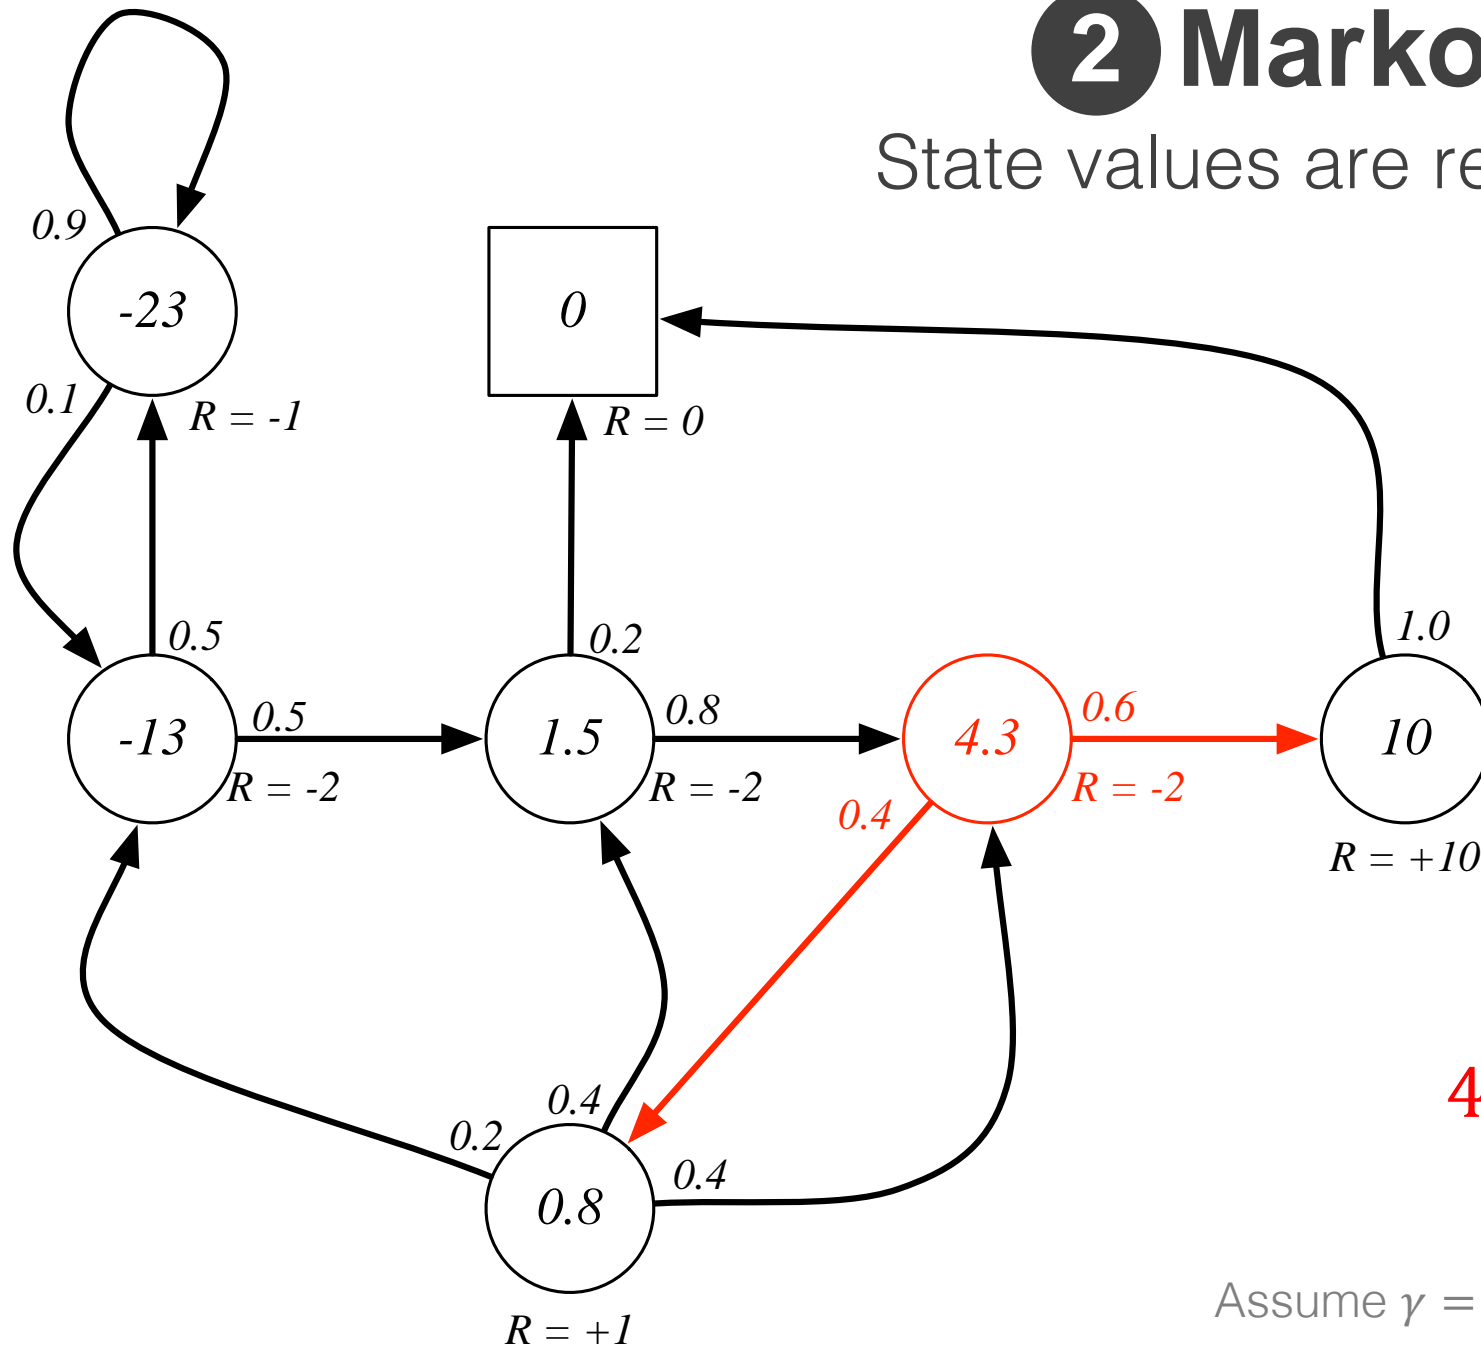
$$= E[R_{t+1} + \gamma v(s_{t+1}) | S = s_t]$$

This recursive relationship is a version of the **Bellman Equation**

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

State values are related to neighboring states



possible states we could transition to from  $s$

$$v(s) = E[R_s + \gamma v(s') | s]$$

$$v(s) = R_s + \gamma \sum_{s'} P_{ss'} v(s')$$

$$4.3 = -2 + 0.6 \times 10 + 0.4 \times 0.8$$

Notation:  $s = s_t$  and  $s' = s_{t+1}$

$R_s = E[R_{t+1} | S_t = s]$

Assume  $\gamma = 1$

Example from David Silver, UCL, 2015




# 3 Markov Decision Process

## Components:

State space  $S$ ,  
Transition probabilities,  $P$   
Rewards,  $R$   
Discount rate,  $\gamma$   
Actions,  $A$

Adds interaction with the environment

An agent in a state chooses an action, the environment (the MDP) provides a reward and the next state



Example from David Silver, UCL, 2015

# 3 Markov Decision Process

Policy (how we choose actions)

(can be stochastic or deterministic)

$$\pi(a|s) = P(a|s)$$

State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t|s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s')|s]$$

Action value function


(expected return from state  $s$ , taking action  $a$ , and following policy  $\pi$ )

$$q_{\pi}(s, a) = E[G_t|s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a')|s, a]$$

$$R_s^a = E[r_{t+1}|S_t = s, A_t = a]$$

Example from David Silver, UCL, 2015




# Bellman Expectation Equations for the **state value** function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t | s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s') | s]$$

$$R_s^a = E[R_{t+1} | S_t = s, A_t = a]$$



Expectation over the possible actions

Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

$$v_{\pi}(s) = \underbrace{\sum_a}_{\text{Expectation over actions}} \underbrace{\pi(a|s)}_{\text{Expectation over rewards}} \left( \underbrace{R_s^a + \gamma \sum_{s'} P_{ss'}^a v_{\pi}(s')}_{\text{Expectation over next possible states}} \right)$$

# Bellman Expectation Equations for the **action value** function


(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )

$$q_{\pi}(s, a) = E[G_t | s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a') | s, a]$$

$$R_s^a = E[R_{t+1} | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) \leftarrow s, a$$



Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

Expectation over the possible actions

$$q_{\pi}(s, a) = \underbrace{R_s^a}_{\text{orange}} + \gamma \underbrace{\sum_{s'} P_{ss'}^a}_{\text{blue}} \underbrace{\sum_{a'} \pi(a' | s') q_{\pi}(s', a')}_{\text{green}}$$


# Bellman Expectation Equations

## State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t | s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s') | s]$$




$$v_{\pi}(s) = \sum_a \pi(a|s) \left( \underbrace{R_s^a}_{\text{orange}} + \gamma \sum_{s'} \underbrace{P_{ss'}^a}_{\text{blue}} v_{\pi}(s') \right)$$

## Action value function

(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )


$$q_{\pi}(s, a) = E[G_t | s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a') | s, a]$$



$$q_{\pi}(s, a) = \underbrace{R_s^a}_{\text{orange}} + \gamma \sum_{s'} \underbrace{P_{ss'}^a}_{\text{blue}} \sum_{a'} \underbrace{\pi(a'|s')}_{\text{green}} q_{\pi}(s', a')$$

# Example



$$v_{\pi}(s) = \sum_a \pi(a|s) \left( \underbrace{R_s^a}_{\text{orange}} + \gamma \underbrace{\sum_{s'} P_{ss'}^a v_{\pi}(s')}_{\text{blue}} \right)$$

$$v_{\pi}(s_1) = \underbrace{(0.5)(5 - 1)}_{q_{\pi}(s_1, a_1)} + \underbrace{(0.5)[(0.2)(1) + (0.8)(-50)]}_{q_{\pi}(s_1, a_2)} + \underbrace{(0.9)(10) + (0.1)(0)}_{q_{\pi}(s, a) = E[R + \gamma v_{\pi}(s')]} \quad \gamma = 1$$

$\underbrace{R_{s_1}^{a_1} v_{\pi}(s_0)} \quad \underbrace{R_{s_1}^{a_2}} \quad \underbrace{P_{12}^{a_2} v_{\pi}(s_2) + P_{13}^{a_2} v_{\pi}(s_3)}$

$$q_{\pi}(s, a) = E[R + \gamma v_{\pi}(s')]$$

# 3 Markov Decision Process

State value function,  $v_{\pi}(s)$

Assumptions:

$$\pi(a|s) = \frac{1}{\text{\# of reachable states}}$$

(randomly select an action)

$$\gamma = 1$$

(no discounting)


Bellman Expectation Equation

$$v(s) = E[R_s^a + \gamma v(s')|s]$$

$$-1.3 \approx 0.5 \times (-1 - 2.3) + 0.5 \times (-2 + 2.7)$$

Assume  $\gamma = 1$

Example from David Silver, UCL, 2015




# 3 Markov Decision Process

Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies


$$v_*(s) = \max_{\pi} v_{\pi}(s)$$



Example from David Silver, UCL, 2015



# 3 Markov Decision Process



Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$


Optimal **action-value** function,  $q_*(s, a)$

Maximum value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

Example from David Silver, UCL, 2015

# 3 Markov Decision Process



Optimal **policy**,  $\pi_*(s)$   
Which action to take at each moment

$$\pi_*(s) = \arg \max_a q_*(s, a)$$

Once we have the optimal value functions, we've "solved" the MDP!

Example from David Silver, UCL, 2015

# Reinforcement Learning Roadmap

Knowledge of **Environment**

## **Perfect knowledge**

Known Markov  
Decision Process



## **No knowledge**

Must learn from  
experience

## Dynamic Programming

What's a Markov Decision Process?  
How do we find optimal policies?

## Monte Carlo Control

How do we estimate our value functions?  
How do we use the value functions to choose actions?